2025 Retrospective

Product Lead — 2025 Retrospective

1) Overall Goal

1.1 What was Product's primary goal for 2025?

Stated intent (what we thought we were doing):

Ship an MVP Policy Writer within 2–4 weeks of design kickoff that clients could use to upload policies and receive a gap analysis.

What this actually was:
A **speed-based delivery plan**, not a product outcome.

1.2 Was that goal explicit and measurable?

Explicit?
Partially.

The goal was explicit. But several factors including development cadence/commitment, migration, security, scope creep, poor product definition, discovering requirements late, HIPAA/SOC, working on MHRS service separately led to increased delivery time.

Measurable?
Weakly measurable.

- The only implied metric was **"did we ship something quickly?"**
- There was **no definition** of:
    - what "online" meant
    - what accuracy meant
    - what "usable" meant
    - what counted as success with real clients

There was **no single binary success condition**.

1.3 Did we achieve it?

Status: Not at all.

Evidence:

- Policy writer:

    - ❌ Not in production
    - ❌ Not tested end-to-end
    - ❌ Not used by real clients
    - ❌ No automated gap analysis that can be trusted
- Customers:
    - ❌ No paying facilities using the policy writer
    - ❌ No completed delivery cycle
- Compliance:
    - ❌ Hipaa not understood early
    - ❌ Risk analysis not started (infra instability)

1.4 Why or why not? (Specific, not circumstantial)

Decision failures

- We optimized for **speed of shipping** instead of **truth (working ai), defensibility, and delivery**
- We did not understand early that:
    - compliance would gate architecture
    - accuracy would gate sales
- We allowed “Mvp” to mean “anything that runs”

Scope failures

- Massive scope creep without a kill rule:
    - editor
    - separate gap analysis
    - Carf plans
- Multiple features were built **before one was fully usable**

Execution failures

- No locked definition of done
- No binary milestones
- No fallback delivery lane
- No delivery ownership model

Compliance discovery

- Hipaa + sud Part 2 implications were discovered late
- This forced re-architecture and invalidated early assumptions

Dependency failures

- Multiple infra moves (Contabo → Aws → Gcp)
- No stable environment long enough to harden the product
- Slow engineering velocity relative to scope

---

2) Process Review

2.1 What are the current core processes in Product?

Actual processes (not idealized):

- Ceo conducts research and customer conversations
- Ceo synthesizes insights mentally
- Ceo is final decision-maker on:
    - what to build
    - when to build
    - whether something ships
- Product documentation:
    - a few videos
    - a few documents
    - heavy reliance on memory
- Build starts quickly with unclear finish criteria
- Scope expands reactively as new ideas emerge

2.2 Which processes consistently worked?

- **Ceo-led customer conversations**
    - Produced real insight
    - identified real problems in the market
- **Fast prototyping**
    - We could build quickly when direction was clear

2.3 Which processes failed or slowed execution?

- **No compliance-first thinking**
    - Led to rework and stalled progress
- **No scope lock**
    - Everything important was treated as urgent
- **No change control**
    - New ideas were added without removing old ones
- **No delivery definition**
    - We didn't ship anything, the definition of delivery kept shifting.
- **No stable execution cadence**
    - Progress was episodic, not compounding

2.4 What workarounds emerged and why?

Workaround 1: Big picture was mostly with Ceo

- Why: no structured synthesis or product system

Workaround 2: Parallel product building

- Why: we had new insights on what the clients would need.

Workaround 3: Manual reasoning instead of tooling

- Why: analyzer accuracy was not trustworthy

Workaround 4: Over-expansion into l.m.s and notetaker

- Why: pressure to show progress despite lack of core readiness

---

3) Measurement and Tracking

3.1 How are we measuring success today?

Reality:

- Informally
- Based on:
    - what feels closer
    - what demos better
    - what unblocks conversation
- No single scorecard
- No weekly outcome tracking

3.2 Which metrics actually influenced decisions?

None consistently.

Decisions were driven by:

- urgency
- intuition
- perceived opportunity
- external pressure

3.3 Which metrics were tracked but not useful?

- We monitored development.

These did **not** correlate with customer value or revenue.

3.4 How do we track progress week to week?

- No consistent system
- No locked weekly outcome
- No falsifiable milestones

Proposed Changes for 2026 (Product-Led)

Goal changes

- One company-level outcome per quarter
- Outcome must include:
    - delivery
    - payment
    - usage

Process changes

- Consultancy-first delivery until product earns trust
- Hybrid delivery allowed but standardized
- Hard milestone ladder (binary, dated)
- Scope creep kill-switch

Metric changes

- Weekly scorecard:
    - facilities in motion
    - delivered + paid
    - cycle time
    - analyzer error rate
- No metrics without decision consequence

Final Truth (for the retreat)

This is not a failure of intelligence or effort.

It is a **discipline gap**:

- unclear goals
- undefined delivery
- unprotected scope
- late compliance awareness

The Q1 2026 plan you now have directly corrects these failures.

2026 Plans

Theraptly — 2026 Product Plan + okrs

1) 2026 Product Thesis

Theraptly earns trust by becoming legally defensible first, operationally repeatable second, and scalable only after controls and evidence are proven.

Sequence: **Hipaa readiness (defined) → consultancy-led delivery + consultant self-serve → ai hardening → facility self-serve → Notetaker scale → soc 2 audit → Billing**

---

2) Product Portfolio (Roles Locked)

1. **Policy Writer / Analyzer**
- Q1: consultancy-led delivery + **consultant self-serve**
- Q2: facility self-serve begins
- Trust + revenue anchor
2. **lms**
- Q1: finish + deploy + upsell
- Training evidence + expansion revenue
3. **Notetaker**
- Q1: build + controlled pilots
- Q2: massive deployment
- Daily-use growth engine
4. **Billing Assistant**
- Q3: build (pilot-only, high risk)
- Q4: validate (reduce rejected claims)
- Weekly-use retention driver

---

3) Definitions (Non-Negotiable)

3.1 “Hipaa compliant” in Q1 means this (and only this)

Hipaa Q1 Readiness Definition (Hq1-r):
Theraptly may claim “Hipaa-ready / Hipaa-aligned operations” only when all are true:

1. **Risk analysis completed** and signed (v1, not draft)

2. **Required Hipaa policies live** (privacy/security/breach + workforce training + sanctions)
3. **Access controls enforced** (role-based access, least privilege, MFA where applicable)
4. **Audit logging live** for Phi-touch events (access/upload/analysis/export/admin changes)
5. **Incident response + breach procedure tested once** (tabletop exercise)
6. **Baa posture defined** for all vendors that touch ePhi (or vendor removed)

Explicit boundary: This is not a certification or government stamp.
Sales-safe wording: "We operate with Hipaa-aligned controls and a completed risk analysis."

3.2 42 Cfr Part 2 (Q1) — what it means operationally

Part 2 readiness in Q1 means:

- Data classification exists for Part 2 programs
- Access restrictions and disclosure rules are enforced
- Staff trained on "extra-sensitive sud data"
- Outputs and notifications avoid sensitive metadata leakage

3.3 soc 2 in 2026 — timing locked

- Q2: audit readiness (controls operating + evidence)
- **Q3: soc 2 Type I audit completed (mock type ii audit).**

---

4) A.i Engine Accuracy (Defined Properly)

4.1 What is a "gap assertion" (unit of accuracy)

A **gap assertion** is one claim of the form:

"Policy set lacks requirement X" or "Requirement X is present but inadequate" or "Requirement X is present."

4.2 What counts as accurate

A gap assertion is accurate only if it is:

- **Quote-anchored** to exact text, or
- Correctly marked **“Not found in uploaded documents”**, and
- Classified correctly (present / inadequate / not found)

4.3 The Q1 accuracy target (99.9% defensible accuracy)

99.9% defensible accuracy =
≥99.9% of gap assertions in the test set are accurate by the definition above.

4.4 Coverage floor (prevents “refuse to answer” gaming)

The engine must produce a usable analysis for:

- **≥95% of uploaded documents** (coverage floor)

4.5 Safety label requirement (until facility self-serve is mature)

All exported reports in Q1 must include:

- “Tool-assisted analysis; final responsibility rests with reviewer”
- “No claim is made without quote anchoring or ‘not found’”

---

5) Q1 Operating Model

Q1 Access Model (crucial)

- **Facilities:** receive deliverables; do **not** self-run analysis yet
- **Consultants:** get **full self-serve access** in Q1 (gated by onboarding + terms)

Why this is the right ladder

Consultants are controlled professional users; facilities are broader and higher-risk.
Q1 focuses on trust, accuracy, and compliance evidence.

---

6) 2026 Outcomes (One per Quarter)

Q1 Outcome — Hipaa readiness + consultant self-serve + paid delivery

By Mar 31, 2026, we operate at Hq1-r (Hipaa readiness defined above), deliver and collect payment for 10 facilities using Theraptly-run analysis + consultant self-serve, ship lms for upsell, and harden the ai engine to 99.9% defensible accuracy.

Q2 Outcome — facility self-serve begins + Notetaker massive push

By end of Q2, facility self-serve policy analysis is live safely, and Notetaker is deployed at scale with daily active usage.

Q3 Outcome — soc 2 audit + Billing build

By end of Q3, soc 2 Type I audit is completed successfully, and Billing Assistant v1 is piloted.

Q4 Outcome — Billing validation + ecosystem expansion

By end of Q4, Billing reduces rejected claims in pilots and multi-product adoption increases revenue per facility.

---

7) Guardrails (Prevents 2025 repeat)

7.1 Consultant self-serve gating (controls added)

Consultants get access only after:

1. 30-minute onboarding (recorded)
2. Terms accepted (use + disclaimers)
3. First analysis reviewed once with pm/qa checklist (lightweight)
4. Output labeling is non-removable in exports

7.2 Facility self-serve (Q2) rollout kill-switch

Facility self-serve cannot roll out if any of these occur in pilot:

- fabricated claim rate > **0%**
- accuracy < **98.5%**
- support burden > **10 tickets per 100 analyses**
- incident involving Phi mishandling

If triggered: rollback to consultant-only self-serve, facilities return to “deliverable-only” temporarily.

7.3 pm load protection (hard rules)

During Q1, pm cannot own everything. So:

- **Design becomes Product Execution Arm**
  - templates, flows, onboarding screens, checklists, export formatting
- **ai engineer owns accuracy work** (pm defines specs + verifies)
- pm focuses on:
  - definitions, gating, quality acceptance, delivery factory, weekly scorecard

Kill-switch for overload:
If ai accuracy is behind plan by >2 weeks, pm pauses new product asks (non-core) and reduces lms expansion tasks to “minimum sellable.”

7.4 Sales-safe compliance language (prevents overclaiming)

Ceo and Growth must use approved phrases:

- ✅ “Hipaa-aligned controls and completed risk analysis”
- ✅ “soc 2 audit planned for Q3”
- ❌ “We are Hipaa certified”
- ❌ “We are fully compliant in every way”

---

8) 2026 okrs

Q1 2026 okrs — Hipaa readiness + ai accuracy + consultant self-serve + lms

Objective 1: Reach Hipaa Q1 readiness (Hq1-R) and Part 2 readiness

Key Results

1. Risk analysis v1 completed and signed.
2. Hipaa + Part 2 workforce training completed by 100% of team.
3. Audit logs live for access/upload/analysis/export/admin changes.
4. Incident response + breach tabletop executed once with action items closed.

Objective 2: Harden ai engine to 99.9% defensible accuracy with coverage

Key Results

1. ≥99.9% accurate gap assertions on a test set of ≥100 real policy documents.
2. Coverage floor met: analysis succeeds on ≥95% of uploaded documents.
3. Fabricated claim rate = **0%**.
4. Manual correction rate ≤10% of outputs by end of Q1.

Objective 3: Deliver paid facility outcomes via Theraptly consultancy-led workflow

Key Results

1. 10 facilities delivered + fully paid (contracts + evidence packs).
2. Median delivery cycle time ≤10 business days (upload complete → delivered).
3. 100% of facilities receive corrected pack + change log + evidence pack.
4. Ceo completes 10 sales cycles without compliance retraction.

Objective 4: Make Consultant self-serve real, and lms sellable

Key Results

1. ≥10 consultants onboarded (placeholders allowed) with gated access completed.
2. 80% MoM repeat usage among active consultants (activity definition locked).
3. lms shipped + used in ≥70% of delivered facilities within 30 days.
4. lms training evidence included in ≥80% of facility evidence packs.

Q2 2026 okrs — Facility self-serve + Notetaker massive push

Objective 1: Launch facility self-serve safely

Key Results

1. Facility self-serve enabled for policy analysis.
2. ≥50% of new analyses run directly by facilities.
3. No degradation: accuracy stays ≥98.5%, fabricated claims stay 0%.

Objective 2: Deploy Notetaker at scale

Key Results

1. ≥1,000 therapists onboarded.
2. ≥70% DAU/WAU.
3. ≥90% of recordings generate SOAP notes.
4. NPS ≥6.

Q3 2026 okrs — soc 2 audit + Billing build

Objective 1: Complete soc 2 Type I audit

Key Results

1. soc 2 Type I audit completed successfully.
2. 0 critical findings; ≤3 minor findings.
3. Remediation plan approved within 30 days.

Objective 2: Build Billing Assistant v1 (pilot-only)

Key Results

1. Screen/context capture works reliably in target workflows.
2. ai suggestions usable in ≥2 billing workflows.
3. Pilot launched in ≥2 facilities.

Q4 2026 okrs — Billing validation + ecosystem expansion

Objective 1: Prove Billing impact

Key Results

1. Rejected claims reduced by ≥20% in pilot facilities.
2. Weekly usage by ≥70% of pilot users.

Objective 2: Increase revenue per facility

Key Results

1. ≥60% of facilities use 2+ products.
2. ≥30% use 3+ products.
3. ARPF up ≥40% vs Q1. (Average Revenue per Facility)

Abbreviations & Term Keys

Core Compliance & Legal

Hipaa
Health Insurance Portability and Accountability Act (U.S.). Governs privacy, security, and breach handling for protected health information (Phi).

Hq1-r *Hipaa Q1 Readiness.*
Theraptly's internal definition of "Hipaa-aligned operations" for Q1, meaning: completed risk analysis, required policies, access controls, audit logs, training, and incident response testing.
Not a certification.

42 Cfr Part 2 (Part 2)
U.S. federal regulation governing substance use disorder (sud) records. Stricter than Hipaa. Limits disclosure, access, and metadata exposure.

soc 2
Service Organization Control 2. Independent audit standard for security, availability, confidentiality, processing integrity, and privacy.

soc 2 Type I
Audit that evaluates whether controls are *designed correctly* at a point in time (not operating effectiveness over time).

BAA
Business Associate Agreement. Contract required when vendors handle ePhi on behalf of a covered entity.

Phi / ePhi
Protected Health Information / electronic Phi.

Product & Usage Metrics

FAU-PW
Facility Active Use – Policy Writer.
A facility counts only after contract, payment, uploads, gap delivery, corrected pack delivery, and evidence completion.

CMR
Consultant Monthly Repeat usage.
Percentage of consultants active in month N who are active again in month N+1.

lms
Learning Management System. Used to deliver staff training tied to compliance and policy requirements.

DAU
Daily Active Users.
For Notetaker: a user who records a session and generates SOAP notes on the same day.

WAU
Weekly Active Users.

NPS
Net Promoter Score. Measures user willingness to recommend the product.

ARPF
Average Revenue Per Facility.

---

ai & Accuracy

ai
Artificial Intelligence.

Gap Assertion
A single analyzable claim made by the Policy Analyzer, e.g.:

- “Requirement X is missing”
- “Requirement X is present but inadequate”
- “Requirement X is present”

Quote-Anchored
An analyzer output that references exact, contiguous text from the uploaded policy document.

Coverage Floor
Minimum percentage of documents for which the analyzer must produce usable output (set at ≥95%).

Fabricated Claim
Any assertion not quote-anchored and not explicitly marked “not found.”
Target: **0% allowed**.

Infrastructure & Engineering

AWS (a.w.s.)
Amazon Web Services.

GCP (g.c.p.)
Google Cloud Platform.

MFA
Multi-Factor Authentication.

SLA
Service Level Agreement (delivery or performance commitment).

API
Application Programming Interface.

Business Model & Go-to-Market

B2B
Business-to-Business. Facilities, organizations, consultants.

B2C
Business-to-Consumer. Individual therapists.

Self-Serve
Users run analyses or workflows directly without Theraptly staff intervention.

Hybrid Delivery
Combination of automated output and human review/correction.

Upsell
Selling additional products (e.g., lms, Notetaker) after the initial Policy Writer engagement.

Roles

pm
Product Manager.

Ceo
Chief Executive Officer.

qa
Quality Assurance.

Reporting & Documentation

Evidence Pack
Required proof set for each facility engagement:

- Contract
- Payment proof
- Upload completion
- Gap report
- Corrected policy pack
- Timestamps and logs

Change Log
Record of all edits made to policy documents during correction.

Planning & Execution

KR
Key Result (used in okrs).

okr
Objectives and Key Results.

Kill-Switch
Pre-defined condition that immediately pauses or rolls back a feature or rollout.

Pilot
Limited, controlled deployment to validate safety and value before scale.

Q1 2026

Theraptly Q1 2026 Retreat Handout

Protected Sprint + Parallel Work + Compliance Gates + Hard Milestones + Timeline + RACI

Quarter window: Jan 4 to Mar 31, 2026
Protected sprint window: Jan 12 to Feb 15, 2026
Sprint focus: Backend Policy Analyzer defensibility + h.i.p.a.a.-ready gate + consultant-operated workspace (consultant self-serve)

---

0) Current truth (constraints)

- Policy Writer is not in production and not tested end-to-end
- Automated gap analysis is not production-grade
- Infra is in-flight: a.w.s. to g.c.p. migration
- Not yet h.i.p.a.a.-ready; formal risk analysis not complete
- h.i.p.a.a. training starts early Q1; p.m records training video this week
- New engineer joins for ai plus policy analyzer work during sprint start
- No consultant workspace yet
- l.m.s is about 70% done; notetaker about 40% done
- Sales model:
    - **Chad** runs most facility calls and pipeline execution
    - **Mr. Walker (Ceo)** takes some facility calls plus partnerships, facility associations, executives, ex-execs, ex carf auditors
- Operating posture for Q1: hybrid delivery is expected (system drafts + human verification + corrected docs)

---

1) Company outcome for Q1 (non-negotiable)

By Mar 31, 2026, we close, deliver, and collect payment for 10 facility consultancy engagements end-to-end, using the Policy Writer workflow as the delivery system, with full evidence packs and no compliance retractions.

---

2) Protected sprint objective (Jan 12 to Feb 15)

By Feb 15, 2026, we have a backend-only Policy Analyzer that is 99.9% defensively accurate, reference-anchored, and trusted enough for Chad and Ceo led consultancy delivery and consultant-operated self-serve use.

This sprint objective gates scale for the rest of Q1.

---

3) What counts (definitions locked)

3.1 Facility Active Use for Policy Writer (f.a.u.-p.w)

A facility counts only when all are true:

1. Contract signed
2. Deposit collected (mandatory)
3. Policy inventory checklist agreed
4. At least 90% uploaded into the document hub
5. Gap report delivered (hybrid allowed)
6. Corrected policy pack delivered
7. Final payment collected
8. Evidence pack complete

No evidence equals it didn't happen.

3.2 Deposit rule (mandatory)

No deposit, no delivery slot. Work does not start until deposit is received.

3.3 Upload deadline rule (mandatory)

Upload due within 5 business days after deposit.
If not complete by Day 5:

- delivery clock pauses
- facility moves to "stalled upload"
- delivery is rescheduled

---

4) Q1 access model (controlled, non-negotiable)

4.1 Consultants: consultant-operated workspace (Q1 enabled)

Consultants can self-serve in Q1, but inside a constrained workspace.

Consultants may:

- upload policies
- run analysis
- export draft gap reports
- generate correction drafts
- manage multiple facilities they are assigned

Consultants may not:

- remove system labels or disclaimers
- represent outputs as certification
- override “not found” without human review trail
- export without m.c.c plus change log attachment

4.2 Facilities: no self-serve in Q1

Facilities receive deliverables in Q1. Facility self-serve begins in Q2.

---

5) What “99.9% accuracy” means (Policy Analyzer)

5.1 Defensive accuracy rules (truthfulness guarantees)

- The analyzer never invents policy content
- The analyzer never asserts “your policy says” unless it can quote it
- Every output item is either:
    - anchored to exact policy text (quote + location), or
    - explicitly “not found in uploaded documents”

5.2 Measurement unit

Accuracy is measured per **gap assertion** (not per document).

5.3 Coverage floor (prevents refusal-gaming)

Analyzer must produce usable output for at least 95% of uploaded documents.

5.4 Mandatory output labeling in Q1

Every export includes:

- “tool-assisted draft, requires reviewer verification”
- “no invented claims, quote anchored or not found”

---

6) Deliverable quality gates (shipping rules)

No deliverable ships unless all three pass:

Gate A: truthfulness gate

- fabricated claims rate = 0
- quote-anchored or not found only

Gate B: coverage gate (m.c.c)

Each facility deliverable includes a Minimum Coverage Checklist with:

- present and adequate
- present but inadequate
- not found

Gate C: change-log gate

Corrected policy pack includes a Correction Change Log:

- what changed
- where (policy and section)
- why (gap addressed)
- date and owner

---

7) H.i.p.a.a.-ready in Q1 (definition + evidence gate)

H.i.p.a.a.-ready by Q1 means all are true:

1. Risk analysis v1 completed and signed
2. Required h.i.p.a.a. policies live
3. workforce training completed by 100% of team
4. role-based access enforced
5. audit logs live for upload, access, analysis generation, export, and admin changes
6. incident response tabletop completed once with action items closed

Sales-safe wording only:

- “h.i.p.a.a.-aligned operations with completed risk analysis” Not allowed:
- “h.i.p.a.a certified”
- “fully compliant in every way”

---

8) New guardrails added

8.1 Accuracy grace band rule (keeps 99.9% without freezing progress)

- Target remains **99.9%**.
- If accuracy is **99.7% to 99.89%**, we do not pretend it is 99.9%.
    - We continue delivery only with mandatory secondary review on every report.
    - We do not expand consultant self-serve volume until 99.9% is achieved.
- If accuracy is **below 99.7%**, protected sprint scope narrows further:
    - stop parallel distractions
    - fix top error patterns only until we return to grace band

8.2 Consultant export traceability (liability hygiene)

Every consultant export must contain:

- consultant id
- facility id
- timestamp
- analyzer version id
- report version id

Audit logs must record:

- who exported
- what was exported
- when it was exported

No export is considered valid without traceability.

8.3 Sales ramp checkpoint (Jan 31)

By **Jan 31**, we must have at least:

- 4 proposals sent
- 2 deposits collected

If missed:

- increase top-of-funnel immediately (more outbound, associations, partnerships)
- Ceo temporarily increases deal involvement
- reduce non-core internal work to protect sales throughput

8.4 l.m.s stop condition (prevents silent failure)

If l.m.s evidence is attached to fewer than 50% of delivered facilities by **Facility 04**, then:

- l.m.s upsell pauses for 2 weeks
- fix onboarding, reporting, and evidence packaging

- resume upsell only after attachment rate recovers

---

9) Strategy (how we win given constraints)

- Consultancy sells now (Chad volume, Ceo strategic trust and partnerships)
- Delivery is hybrid in Q1
- Product stabilizes while delivering (stability beats features)
- Compliance posture advances weekly, not later
- Migration cannot pause sales or delivery (fallback lane exists)
- l.m.s is completed and upsold in Q1 using minimum sellable definition
- Notetaker continues build in Q1, massive push in Q2
- s.o.c 2 audit is Q3, prep begins after Q1 stabilization

---

10) What we will not do in Q1 (scope lock)

- No facility self-serve policy analysis
- No UI polish beyond functional basics
- No “certification” claims (h.i.p.a.a or s.o.c 2)
- No s.o.c 2 audit in Q1 (prep only)
- No billing assistant work
- No new major product expansions that do not unblock the 10-facility outcome

---

11) Workstreams and owners (Q1)

A) Consultancy delivery engine (Owner: p.m)

- delivery factory
- m.c.c, change log, evidence pack discipline
- quality gates enforcement

B) Policy Analyzer backend (Owner: ai engineer, accountable: eng lead)

- parse, anchor, export, versioning
- no guessing, no “helpful” speculation

C) h.i.p.a.a readiness (Owner: compliance owner, supported by eng lead and p.m)

- training, policies, risk analysis, logs, tabletop

D) Sales and partnerships (Owner: Chad for pipeline, Ceo for strategic)

- calls, proposals, deposits

- association partnerships, auditors, exec networks

E) l.m.s completion and upsell (Owner: l.m.s owner, support: design and p.m for evidence packaging)

- minimum sellable definition
- attach evidence to facility packs

F) Notetaker (Owner: notetaker owner)

- build progress in Q1
- no massive deployment until Q2

---

12) Milestone ladder

M0: Retreat lock (Jan 12)

- f.a.u.-p.w, deposit and upload rules locked
- m.c.c categories locked
- change-log template locked
- weekly scorecard locked
- consultant export traceability spec approved
- scope lock signed

M1: g.c.p staging stable (Jan 15)

- staging deploy exists
- upload works (success and failure paths)
- export works (basic ok)

M2: h.i.p.a.a training kickoff (Jan 15)

- training video recorded
- training schedule published
- first quiz completed

M3: Analyzer core loop working (Jan 19)

- upload to parse to extract to export runs end-to-end
- missing data produces “not found”

M4: Gap draft generation (Jan 26)

- anchored draft gaps generated
- p.m can produce a usable hybrid gap report

M5: 99.9% accuracy proven (Feb 2)

- test set completed on real policies
- fabricated claims = 0
- coverage floor met

M6: Consultancy simulation passes (Feb 9)

- full mock delivery done
- m.c.c, change log, evidence pack produced
- traceability included in exports

M7: Sales readiness gate (Feb 15)

- Chad and Ceo safe promises aligned
- demo outputs package ready
- consultant workspace enabled under constraints

M8: h.i.p.a.a-ready achieved (Feb 21)

- risk analysis v1 signed
- audit logs live
- tabletop done

M9: 3 facilities delivered and paid (Feb 28)

- f.a.u.-p.w complete for Facilities 01 to 03

M10: 8 facilities delivered and paid (Mar 21)

M11: 10 facilities delivered and paid (Mar 31)

---

13) Fallback lane (migration cannot pause delivery)

If g.c.p staging blocks delivery, delivery continues via one approved secure interim lane.

Rules:

- only one approved channel, no ad-hoc sharing
- access limited to named roles
- every file named with facility id, doc type, version
- logged in evidence pack
- imported into document hub once restored
- removed from interim channel after import

---

14) Weekly cadence

Monday: targets, facility stage assignments, promise review (Ceo, Chad, p.m)
Tuesday: closing day (walkthroughs, deposits)
Wednesday: delivery day (p.m production, a.i fixes top error pattern)
Thursday: trust and payment day (delivery walkthroughs, final payment, referrals)
Friday: ship and retro (stability fixes, post-mortems)

---

15) Timeline (Jan 5 to Mar 31)

Jan 5 to Jan 12 (prep week)

- corporate retreat
- record l.m.s training video
- lock analyzer “never do” rules
- assemble policy corpus for testing
- lock m.c.c categories and change-log template
- publish sprint calendar

Jan 12 to Feb 15 (protected sprint)

Week 1: onboard ai engineer, ship core loop, daily “does it ever guess” review
Week 2: gap draft generation, internal alpha runs, top error patterns list
Week 3: accuracy hardening, reach 99.9% defensible accuracy
Week 4: consultancy simulation, evidence pack, traceability validation
Week 5: sales readiness, demo package, consultant workspace enabled under constraints

Feb 16 to Mar 31 (scale delivery to 10)

- Feb 16 to Feb 28: Facilities 01 to 03 delivered and paid
- Mar 1 to Mar 21: Facilities 04 to 08 delivered and paid, l.m.s stop condition monitored
- Mar 22 to Mar 31: Facilities 09 to 10 delivered and paid, Q1 retro, Q2 facility self-serve planning

---

16) RACI table (final)

Roles:

- **Ceo:** Mr. Walker
- **sales:** Chad
- **p.m:** Product Manager (you)
- **ai:** ai/policy engineer

- **eng:** engineering lead
- **des:** design execution arm
- **l.m.s:** l.m.s owner
- **note:** notetaker owner
- **comp:** compliance owner (pm)

Work item	Ceo	sales	p.m	ai	eng	des	l.m.s	note	comp
Q1 outcome governance	A	R	R	I	I	I	I	I	R
Define 99.9% accuracy rules	C	I	A/R	R	C	C	I	I	R
Analyzer backend build	I	I	C	R	A	I	I	I	C
Export traceability enforcement	I	I	A/R	C	R	I	I	I	C
Consultant workspace scope rules	C	I	A/R	C	I	C	I	I	R
Facility sales pipeline	C	A/R	I	I	I	I	I	I	I
Partnerships and associations	A/R	C	I	I	I	I	I	I	I
Facility intake, deposit enforcement	I	A/R	R	I	I	C	I	I	C
Hybrid gap report delivery	I	I	A/R	C	I	C	I	I	C
Corrected policy pack production	I	I	A/R	C	I	C	I	I	C
m.c.c and change-log templates	I	I	A/R	C	I	R	I	I	C
Evidence pack discipline	I	I	A/R	I	C	R	C	I	R

g.c.p migration execution	I	I	C	I	A/R	I	I	I	C
h.i.p.a.a training and tracking	I	I	R	I	C	C	I	I	A/R
Risk analysis v1 and tabletop	I	I	R	C	C	I	I	I	A/R
l.m.s minimum sellable completion	I	C	C	I	I	C	A/R	I	I
l.m.s evidence attachment	I	C	R	I	I	C	A/R	I	C
Notetaker Q1 build (no scale)	I	I	I	I	I	C	I	A/R	I

Legend: **R** responsible, **A** accountable, **C** consulted, **I** informed